

# Introducing Malloc

Dr Andrew Scott  
a.scott@lancaster.ac.uk

---

---

---

---

---

---

---

1

## malloc( )

- Provides dynamic memory allocation
  - Allocates memory from process heap
- Returns pointer to requested amount of memory
- Pointer remains valid until call to *free( )*
  - No automatic garbage collection -- DIY
- *malloc( )* must be followed by one, and only one, *free( )*

---

---

---

---

---

---

---

2

## In use...

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

char *
first( ) {
    char * first_buff = ( char * ) malloc( 7 );

    return strcpy( first_buff, "Hello" );
}

char *
second( ) {
    char * second_buff = ( char * ) malloc( 7 );

    return strcpy( second_buff, "World!" );
}

int
main( ) {
    char * str1 = first( );
    char * str2 = second( );

    printf( "%s %s\n", str1, str2 );

    free( str1 ); free( str2 );
}
```

---

---

---

---

---

---

---

3

A simple `malloc()`: Structure



---

---

---

---


---

---

---

4

Simple `malloc()`: allocating



- Simple approach...
  - Chain all blocks together: *Previous and Next pointers*
  - Search chain/ list for free space when needed
    - Maintain *Free/ Used* indicator

---

---

---

---


---

---

---

5

Simple `malloc()`: allocating II



---

---

---

---

---


---

---

6

### Simple `malloc()`: the user view...

- `malloc()` returns start address of free area
- All within one process space – no protection
  - User must never write outside returned memory
    - Would overwrite metadata and break lists
    - ...but it's their data they'd be corrupting



7

---

---

---

---

---


---

---

---

### Simple `malloc()`: freeing space

- Assume two used blocks



- `free( myptr )` : Coalesce adjacent free blocks



8

---

---

---

---

---

---

---

---

### Heap Allocation Test: The real `malloc()` and `free()`

```
#include <stdio.h>
#include <stdlib.h>

int
main() {
    void *  addr1, addr2, addr3, addr4, addr5, addr6;

    addr1 = malloc( 1024 );  addr2 = malloc( 1024 );
    free( addr1 );  free( addr2 );

    addr3 = malloc( 1024 );  addr4 = malloc( 1024 );
    free( addr3 );  free( addr4 );

    addr5 = malloc( 512 );  addr6 = malloc( 512 );
    free( addr5 );  free( addr6 );

    printf( "Addr1 = %p, Addr2 = %p\n", addr1, addr2 );
    printf( "Addr3 = %p, Addr4 = %p\n", addr3, addr4 );
    printf( "Addr5 = %p, Addr6 = %p\n", addr5, addr6 );
}
```

9

---

---

---

---

---

---

---

---



### malloc( ) Free Bins

**Note:** The free areas can be anywhere within heap, and can be in any order.

Each bin holds list of free areas  $n$  double words in size, so bin 3 holds areas able to hold  $3 \times 8$ , or 24 bytes -- but these will additionally have  $2 \times 4$  byte LEN/F fields occupying 28 bytes in total

128 Bins in total

---

---

---

---

---

---

---

---

---

---

13

### Growing the Heap: sbrk( )

- When `malloc( )` runs out of space it calls `sbrk( )`
- Heap has *high water mark*, the program *break*
  - Read using: `sbrk( 0 )`
- Can be pushed up `sbrk( +ve )` or pulled down `sbrk( -ve )`
  - In practice, generally snapped to `page* boundary`
- Can be set explicitly using `brk( addr )`

\* As we'll see later, hardware chunks memory into (typically) 4K blocks called *pages*

---

---

---

---

---

---

---

---

---

---

14

### Allocating whole pages: mmap( )

- For large requests, whole *pages* allocated in separate memory region using `mmap( )`
- More efficient and avoids large void on heap seen if two small allocations bookend a free multi-page area

```
/*
 * Allocate single page, at any/ best (NULL) address,
 * PRIVATE (non-shared)
 * and ANONYMOUS, so not backed by a file
 */
char * p = mmap( NULL, sysconf( _SC_PAGE_SIZE ),
                 PROT_READ | PROT_WRITE,
                 MAP_PRIVATE | MAP_ANONYMOUS,
                 -1, 0
               );
```

---

---

---

---

---

---

---

---

---

---

15

## *mmap()*, full example...

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/mman.h>
#include <unistd.h>

int
main() {
    char * p = mmap( NULL, sysconf( _SC_PAGE_SIZE ),
                     PROT_READ | PROT_WRITE,
                     MAP_PRIVATE | MAP_ANONYMOUS,
                     -1, 0
                    );

    if( p <= 0 ) {
        perror( "mmap" );
        exit( EXIT_FAILURE );
    }

    strcpy( p, "Hello World!" );
    puts( p );
}
```

16

## If *sbrk* adds memory, why *malloc*?

- *sbrk()* returns or adjusts high water mark
  - No knowledge of what is below this mark
- *malloc()* tracks/ manages free space
  - Knows which areas free and which are used
  - Splits free areas down to required size
- Beware: Don't forget that *malloc()* uses *sbrk()*
  - *malloc()* can't assume it's the only thing using heap memory
  - Neither can you!
    - For  $p = \text{sbrk}(n)$  you can only use byte addresses  $p..p+n-1$ ,  $n > 0$

17

## *malloc()* Final Thoughts

- Number of alternative approaches, such as...
  - Tree structures, and Buddy algorithm that we'll see later
- Three main problems
  - How to limit heap size by avoiding unnecessary free space
    - i.e. how to keep *sbrk()* high water mark as low as possible
  - How to limit, or manage, memory fragmentation
    - Small amounts of unused space between allocated memory
    - Note: we can't shuffle things in memory as we don't know how programs are using memory
  - How to choose between two or more free areas
    - And should newly freed areas be at head or tail of free list

18